

The kl inequality

Yevgeny Seldin

Target

- Derive an inequality that is tighter than Hoeffding's

Basics in Information Theory


- Entropy of a distribution p

$$H(p) = - \sum_x p(x) \ln p(x)$$

- Binary entropy

- Entropy of a Bernoulli distribution $(1 - p, p)$

$$H(p) = \underbrace{-(1 - p) \ln(1 - p)}_{x=0} \underbrace{- p \ln p}_{x=1}$$



The method of types

0 0 0 0

1 0 0 0

1 1 0 0

0 1 1 1

1 1 1 1

0 1 0 0

1 0 1 0

1 0 1 1

0 0 1 0

1 0 0 1

1 1 0 1

0 0 0 1

0 1 1 0

1 1 1 0

0 1 0 1

0 0 1 1

$$\binom{4}{0} = 1$$

$$\binom{4}{1} = 4$$

$$\binom{4}{2} = 6$$

$$\binom{4}{3} = 4$$

$$\binom{4}{4} = 1$$

- All sequences within the same type have the same probability
- The probability of a type is the number of sequences times the probability of an individual sequence
- The probability of observing (the empirical error) $\hat{p}_n = \frac{k}{n}$ is the probability of observing the type $\frac{k}{n}$

$$\mathbb{P}\left(\hat{p}_n = \frac{k}{n}\right) = \binom{n}{k} p^k (1-p)^{n-k}$$

$$\mathbb{P}\left(\hat{p}_n = \frac{k}{n}\right) = \binom{n}{k} p^k (1-p)^{n-k}$$

Bound on the binomial coefficients

- Lemma: for $1 \leq k \leq n-1$

$$\frac{1}{2} \sqrt{\frac{n}{2k(n-k)}} \leq \binom{n}{k} e^{-nH\left(\frac{k}{n}\right)} \leq \frac{1}{2} \sqrt{\frac{n}{k(n-k)}}$$

- Proof:

- $\binom{n}{k} = \frac{n!}{k!(n-k)!}$

- Stirling's approximation of the factorial:

- $\sqrt{2\pi n} \left(\frac{n}{e}\right)^n \leq n! \leq \sqrt{2\pi n} \left(\frac{n}{e}\right)^n e^{\frac{1}{12n}}$

- Three lines of technical derivation (see lecture notes)

- Message: $e^{nH\left(\frac{k}{n}\right)}$ is a good approximation of $\binom{n}{k}$ and $\binom{n}{k} e^{-nH\left(\frac{k}{n}\right)}$ is “small”

- Note: for $k = 0$ and $k = n$ we have $\binom{n}{k} e^{-nH\left(\frac{k}{n}\right)} = 1$, so the message is also valid

$$H(p) = - \sum_x p(x) \ln p(x)$$

Kullback-Leibler (KL) divergence / relative entropy

- “Distance” between probability distributions p and q

- $KL(p||q) = \sum_x p(x) \ln \frac{p(x)}{q(x)} = \mathbb{E}_{X \sim p} \left[\ln \frac{p(X)}{q(X)} \right] = \mathbb{E}_{X \sim p} \left[\ln \frac{1}{q(X)} \right] - H(p)$

- Properties:

- $KL(p||p) = 0$
- $KL(p||q)$ is convex in the pair (p, q)
 - $KL(\lambda p_1 + (1 - \lambda)p_2 || \lambda q_1 + (1 - \lambda)q_2) \leq \lambda KL(p_1 || q_1) + (1 - \lambda)KL(p_2 || q_2)$
- Asymmetry: $KL(p||q) \neq KL(q||p)$

- Binary kl:

$$kl(p||q) = KL((1 - p, p) || (1 - q, q)) = (1 - p) \ln \frac{1 - p}{1 - q} + p \ln \frac{p}{q}$$

$e^{-n \text{kl}\left(\frac{k}{n} || p\right)}$ governs the probability of type $\frac{k}{n}$

$$\begin{aligned} \mathbb{P}\left(\hat{p}_n = \frac{k}{n}\right) &= \binom{n}{k} p^k (1-p)^{n-k} = \binom{n}{k} e^{n\left(\frac{k}{n} \ln p + \frac{n-k}{n} \ln(1-p)\right)} \\ &= \binom{n}{k} e^{-nH\left(\frac{k}{n}\right)} e^{nH\left(\frac{k}{n}\right)} e^{n\left(\frac{k}{n} \ln p + \frac{n-k}{n} \ln(1-p)\right)} \\ &= \underbrace{\binom{n}{k} e^{-nH\left(\frac{k}{n}\right)}}_{\text{"small"}} e^{-n \text{kl}\left(\frac{k}{n} || p\right)} \end{aligned}$$

$$\text{KL}(p || q) = \mathbb{E}_{X \sim p} \left[\ln \frac{1}{q(X)} \right] - H(p)$$

$$\frac{1}{2} \sqrt{\frac{n}{2k(n-k)}} \leq \binom{n}{k} e^{-nH\left(\frac{k}{n}\right)} \leq \frac{1}{2} \sqrt{\frac{n}{k(n-k)}}$$

- Message:

- $\mathbb{P}\left(\hat{p}_n = \frac{k}{n}\right) \approx e^{-n \text{kl}\left(\frac{k}{n} || p\right)}$
- $e^{-n \text{kl}\left(\frac{k}{n} || p\right)}$ governs the probability of observing type $\frac{k}{n}$ when sampling from p

$$\mathbb{E}[X] = \sum_{x \in \mathcal{X}} x \mathbb{P}(X = x)$$

$$\hat{p}_n \in \left\{ 0, \frac{1}{n}, \frac{2}{n}, \dots, \frac{n}{n} \right\}$$

$$\mathbb{P}\left(\hat{p}_n = \frac{k}{n}\right) = \underbrace{\binom{n}{k} e^{-nH\left(\frac{k}{n}\right)}}_{\text{"small"}} e^{-n\text{kl}\left(\frac{k}{n}||p\right)}$$

The kl lemma

- Lemma: $\mathbb{E}\left[e^{n\text{kl}(\hat{p}_n||p)}\right] \leq 2\sqrt{n}$

- Proof:

$$\mathbb{E}\left[e^{n\text{kl}(\hat{p}_n||p)}\right]$$

$$= \sum_{k=0}^n \mathbb{P}\left(\hat{p}_n = \frac{k}{n}\right) e^{n\text{kl}\left(\frac{k}{n}||p\right)}$$

$$= \sum_{k=0}^n \underbrace{\binom{n}{k} e^{-nH\left(\frac{k}{n}\right)}}_{\text{"small"}} \underbrace{e^{-n\text{kl}\left(\frac{k}{n}||p\right)} e^{n\text{kl}\left(\frac{k}{n}||p\right)}}_{=1}$$

$$\leq 2\sqrt{n}$$

The kl lemma is tight

- Lemma: $\mathbb{E}[e^{n\text{kl}(\hat{p}_n||p)}] \leq 2\sqrt{n}$
- Lemma: for $p \in (0,1)$: $\mathbb{E}[e^{n\text{kl}(\hat{p}_n||p)}] \geq \sqrt{n}$
- Proof:

$$\begin{aligned}\mathbb{E}[e^{n\text{kl}(\hat{p}_n||p)}] &= \sum_{k=0}^n \mathbb{P}\left(\hat{p}_n = \frac{k}{n}\right) e^{n\text{kl}\left(\frac{k}{n}||p\right)} \\ &= \sum_{k=0}^n \underbrace{\binom{n}{k} e^{-nH\left(\frac{k}{n}\right)}}_{\text{"small"}} \underbrace{e^{-n\text{kl}\left(\frac{k}{n}||p\right)} e^{n\text{kl}\left(\frac{k}{n}||p\right)}}_{=1} \\ &\geq \sqrt{n}\end{aligned}$$

Markov:

$$\mathbb{P}(X \geq \varepsilon) \leq \frac{\mathbb{E}[X]}{\varepsilon}$$

The kl lemma:

$$\mathbb{E}[e^{n \text{kl}(\hat{p}_n || p)}] \leq 2\sqrt{n}$$

The kl inequality via the kl lemma

- Theorem: $\mathbb{P}\left(\text{kl}(\hat{p}_n || p) \geq \frac{\ln \frac{2\sqrt{n}}{\delta}}{n}\right) \leq \delta$

- Proof

$$\mathbb{P}\left(\text{kl}(\hat{p}_n || p) \geq \frac{\ln \frac{2\sqrt{n}}{\delta}}{n}\right) = \mathbb{P}\left(n \text{kl}(\hat{p}_n || p) \geq \ln \frac{2\sqrt{n}}{\delta}\right)$$

$$\stackrel{\text{Chernoff's bounding technique}}{=} \mathbb{P}\left(e^{n \text{kl}(\hat{p}_n || p)} \geq \frac{2\sqrt{n}}{\delta}\right)$$

Chernoff's
bounding
technique

$$\stackrel{\text{Markov}}{\leq} \frac{\delta}{2\sqrt{n}} \mathbb{E}[e^{n \text{kl}(\hat{p}_n || p)}]$$

$$\stackrel{\text{kl lemma}}{\leq} \delta$$


The kl inequality

- Theorem: $\mathbb{P} \left(\text{kl}(\hat{p}_n || p) \geq \frac{\ln \frac{1}{\delta}}{n} \right) \leq \delta$
- Earlier: $\mathbb{P} \left(\text{kl}(\hat{p}_n || p) \geq \frac{\ln \frac{2\sqrt{n}}{\delta}}{n} \right) \leq \delta$
- Proof:
 - Based on direct derivation (not via the kl lemma); omitted
- The direct derivation is incompatible with PAC-Bayesian analysis
 - There we will need to go via the kl lemma and pay $\ln 2\sqrt{n}$

Relaxations & comparison to Hoeffding

- The kl inequality: $\mathbb{P}\left(\text{kl}(\hat{p}_n||p) \leq \frac{\ln \frac{1}{\delta}}{n}\right) \geq 1 - \delta$
- Pinsker's inequality: $\text{kl}(\hat{p}_n||p) \geq 2(p - \hat{p}_n)^2$
- Corollary: $\mathbb{P}\left(p \leq \hat{p}_n + \sqrt{\frac{\ln \frac{1}{\delta}}{2n}}\right) \geq 1 - \delta$
- Refined Pinsker's inequality: for $p > \hat{p}_n$, $\text{kl}(\hat{p}_n||p) \geq \frac{(p - \hat{p}_n)^2}{2p}$
- Corollary: $\mathbb{P}\left(p \leq \hat{p}_n + \underbrace{\sqrt{\frac{2\hat{p}_n \ln \frac{n+1}{\delta}}{n}}}_{\rightarrow 0 \text{ for } \hat{p}_n \rightarrow 0} + \frac{2 \ln \frac{n+1}{\delta}}{n}\right) \geq 1 - \delta$
 - “Fast convergence rates” (at the rate of $\frac{1}{n}$ rather than $\frac{1}{\sqrt{n}}$)
 - (Significantly) tighter bound for $\hat{p}_n \ll \frac{1}{8}$
 - The kl inequality is even tighter
- Hoeffding:
 - $\mathbb{P}\left(p \leq \hat{p}_n + \sqrt{\frac{\ln \frac{1}{\delta}}{2n}}\right) \geq 1 - \delta$

Inversion of kl



- $\mathbb{P} \left(\text{kl}(\hat{p}_n || p) \leq \frac{\ln \frac{1}{\delta}}{n} \right) \geq 1 - \delta$

- Corollary:


$$\mathbb{P} \left(\text{kl}^{-1-} \left(\hat{p}_n, \frac{\ln \frac{1}{\delta}}{n} \right) \leq p \leq \text{kl}^{-1+} \left(\hat{p}_n, \frac{\ln \frac{1}{\delta}}{n} \right) \right) \geq 1 - \delta$$

$$\text{kl}^{-1+}(\hat{p}_n, \varepsilon) = \max\{p: \text{kl}(\hat{p}_n || p) \leq \varepsilon\}; \quad \text{kl}^{-1-}(\hat{p}_n, \varepsilon) = \min\{p: \text{kl}(\hat{p}_n || p) \leq \varepsilon\}$$

- Inversion of kl:

- $\text{kl}(\hat{p}_n || p)$ is convex in p
- $\text{kl}(\hat{p}_n || \hat{p}_n) = 0$ is the minimum
- $p \in [0, 1]$
- Use binary search on each side of \hat{p}_n


Summary



- kl lemma: $\mathbb{E}[e^{nkl(\hat{p}_n || p)}] \leq 2\sqrt{n}$
- kl inequality: $\mathbb{P}\left(kl(\hat{p}_n || p) \leq \frac{\ln \frac{1}{\delta}}{n}\right) \geq 1 - \delta$
- Pinsker's relaxation: $\mathbb{P}\left(p \leq \hat{p}_n + \sqrt{\frac{\ln \frac{1}{\delta}}{2n}}\right) \geq 1 - \delta$
- Refined Pinsker's relaxation: $\mathbb{P}\left(p \leq \hat{p}_n + \sqrt{\frac{2\hat{p}_n \ln \frac{1}{\delta}}{n}} + \frac{2 \ln \frac{1}{\delta}}{n}\right) \geq 1 - \delta$
 - "Fast rate"
- Direct inversion: $\mathbb{P}\left(kl^{-1-}\left(\hat{p}_n, \frac{\ln \frac{1}{\delta}}{n}\right) \leq p \leq kl^{-1+}\left(\hat{p}_n, \frac{\ln \frac{1}{\delta}}{n}\right)\right) \geq 1 - \delta$
 - Use binary search

Split-kl inequality

- Motivation: the kl inequality is “blind” to the variance
 - $\mathbb{P}\left(\text{kl}(\hat{p}_n || p) \leq \frac{\ln \frac{1}{\delta}}{n}\right) \geq 1 - \delta$



Split-kl inequality

- Solution for discrete random variables $X \in \{b_0, b_1, \dots, b_K\}$:

- Representation as a superposition of Bernoulli random variables

- $\alpha_j = b_j - b_{j-1}$

- $X_{|j} = \mathbb{I}(X \geq b_j)$

- “progress bar”

- $X_{|j}$ is Bernoulli

- $X = b_0 + \sum_{j=1}^K \alpha_j X_{|j}$

- For X_1, \dots, X_n let $\hat{p}_{|j} = \frac{1}{n} \sum_{i=1}^n X_{i|j}$

- $\mathbb{E}[X] = p = b_0 + \sum_{j=1}^K \alpha_j \mathbb{E}[X_{|j}] = b_0 + \sum_{j=1}^K \alpha_j p_{|j}$

- Apply kl inequality to bound the distance between $\hat{p}_{|j}$ and $p_{|j}$ for all j and take a union bound

- Details in the lecture notes

